

Converting to and from Dilated Integers

Rajeev Raman and
David S. Wise, *Member,*
IEEE Computer Society

Abstract—Dilated integers form an ordered group of the Cartesian indices into a d -dimensional array represented in the Morton order. Efficient implementations of its operations can be found elsewhere. This paper offers efficient casting (type) conversions to and from an ordinary integer representation. As the Morton order representation for 2D and 3D arrays attracts more users because of its excellent block locality, the efficiency of these conversions becomes important. They are essential for programmers who would use Cartesian indexing there. Two algorithms for each casting conversion are presented here, including to-and-from dilated integers for both $d = 2$ and $d = 3$. They fall into two families. One family uses newly compact table lookup, so the cache capacity is better preserved. The other generalizes better to all d , using processor-local arithmetic that is newly presented as abstract d -ary and $(d - 1)$ -ary recurrences. Test results for two and three dimensions generally favor the former.

Index Terms—Analysis, programming techniques, numerical algorithms, design, performance, memory hierarchy, caches, paging, compilers, classes, quadrees, octrees, dilation, contraction.

CCS Categories and subject descriptors: E.1 [Data Structures]: Arrays; D.1.0 [Programming Techniques]: General; B.3.2 [Memory Structures]: Design Styles; F.2.1 [Analysis of Algorithms and Problem Complexity]: Numerical algorithms and problems—computations on matrices.

General Term: Algorithms, Design, Performance

Additional Key Words: memory hierarchy, caches, paging, compilers, classes, quadrees, octrees, dilation, contraction.

1 INTRODUCTION

Dilated integers form an ordered group for the Cartesian indices into a d -dimensional array represented in Morton order [1]. Fig. 1 illustrates the Morton order and the use of dilated integers as its Cartesian indices. With an extra high-order bit, they also suffice for casts from Cartesian to/from Ahnentafel indices [2]. The amount of dilation, called d -dilation, varies with the dimension d of the represented array.

Morton order is used to index arrays in applications from computer vision and graphics, cartography and geographical information systems, volumetric analyses such as axial tomography (CAT scans), and the conventional matrix computations of scientific computing. An array in the Morton order can conveniently be decomposed as a 2^d -ary tree whose subtrees have contiguous addresses. Such a locality minimizes page and cache misses. Practically, $d = 2$ for matrices and quadrees and $d = 3$ for 3D aggregates and octrees.

A comfortable programming paradigm to take advantage of such locality uses recursion on 2^d blocks of each array/operand to descend those trees [3], [4]. When the base case is reached (well above 1×1 , typically a 32×32 block), all of it resides in cache.

- R. Raman is with the Department of Computer Science, University Road, Leicester LEQ 7RH, UK. E-mail: rr29@leicester.ac.uk.
- D.S. Wise is with the Computer Science Department, Indiana University, 215 Lindley Hall, Bloomington, IN 47405-7104. E-mail: dswise@cs.indiana.edu.

Manuscript received 29 Dec. 2006; revised 5 July 2007; accepted 1 Aug. 2007; published online 5 Sept. 2007.

Recommended for acceptance by F. Dehne.

For information on obtaining reprints of this article, please send e-mail to: tc@computer.org, and reference IEEECS Log Number TC-0491-1206.

Digital Object Identifier no. 10.1109/TC.2007.70814.




Fig. 1. Morton-order indexing of a 16×16 matrix. As this block-recursive \mathbb{N} order extends across the plane, the dilation of integer 13 indexes the entire 14th row and the doubled dilation of integer 14 indexes the whole 15th column. The transposed Morton-Z order is also available [13].

More importantly, its larger enveloping blocks fit into L2, L3, RAM, swapping disk, and so forth up the memory hierarchy to include the communication of blocks among both shared memory and distributed coprocessors. The programmer need not worry about which size lands where. This implicit mapping of memory has been called “cache oblivious” because she can ignore the details of block size, cache lines, and memory capacity [5]; all of these are local minutiae anyway. Whichever blocks fit will reside in the L2 cache for a good while, with subblocks rolling into L1 repeatedly. Meanwhile, translation lookaside buffer (TLB) misses seem to disappear [6]. This style enhances the portability of code that no longer needs to be retuned to a specific memory architecture.

As these array indexings become better used for their high performance, particularly in the context of cache-oblivious matrix processing, their conversions to and from ordinary integer representations become important. They are essential manifestations, for instance, of the common abstractions from generic programming [7]. Like the I/O conversions of the IEEE floating-point numbers, however, their use really should be infrequent and should be displaced by efficient computation within that type. Nevertheless, unlike the decimal/float conversions [8], [9], knowledge of the details of these algorithms now seems a prerequisite for the acceptance of the type into common use, particularly in the face of lore identifying that conversion as an obstacle [5, p. 288], [10, Section 4.1], [11, Section 4.1].

In order to hasten its promulgation, we offer eight algorithms, two conversions in each direction for $d = 2$ and $d = 3$, to suit specific architectures. In all cases, the conversion is fast, faster than a random access into an array, which is their common runtime context. On current processors, the table lookup is shown to be the faster for repeated castings. For sequential access, however, efficient implementations of the additive operators are directly available on both dilated integers and their cousins, masked integers [12], [13]. Thus, C++ iterators can be even faster than inline conversions for Cartesian indexing into Morton-order matrices.

Section 3 uses table lookup to solve these problems. Section 4 uses register-local operations. Section 4.4 uses an arbitrary d as d -ary and $(d - 1)$ -ary recurrences and Section 5 describes our timings and offers conclusions. New contributions are the halving of tables necessary for contractions in Section 3, new multiplication-based algorithms in Section 4.4, the generalized d -ary and $(d - 1)$ -ary recurrences in Section 4.4, and the demonstration that table lookup is faster on current RISC processors in Section 5.

2 DEFINITIONS

Informally, 2-dilating an integer interleaves one 0 bit between all of the meaningful bits in the binary representation of that integer. As a simple example, consider $i = 2^8 - 1 = 11111111_2 = \text{FF}_{16}$. Its 2-dilation, denoted \vec{i} and read “ i dilated,” is $01010101010101_2 = 5555_{16}$, written $0x5555$ in the C language. Similarly, if j is $0xF0$, then \vec{j} is $0x5500$ and,

```

const unsigned short int dilate_tab2[256] = {
    0x0000, 0x0001, 0x0004, 0x0005, 0x0010, 0x0011, 0x0014, 0x0015,
    0x0040, 0x0041, 0x0044, 0x0045, 0x0050, 0x0051, 0x0054, 0x0055,
    0x0100, 0x0101, 0x0104, 0x0105, 0x0110, 0x0111, 0x0114, 0x0115,
    0x0140, 0x0141, 0x0144, 0x0145, 0x0150, 0x0151, 0x0154, 0x0155,
    0x0400, 0x0401, 0x0404, 0x0405, 0x0410, 0x0411, 0x0414, 0x0415,
    0x0440, 0x0441, 0x0444, 0x0445, 0x0450, 0x0451, 0x0454, 0x0455,
    0x0500, 0x0501, 0x0504, 0x0505, 0x0510, 0x0511, 0x0514, 0x0515,
    0x0540, 0x0541, 0x0544, 0x0545, 0x0550, 0x0551, 0x0554, 0x0555,
    0x1000, 0x1001, 0x1004, 0x1005, 0x1010, 0x1011, 0x1014, 0x1015,
    0x1040, 0x1041, 0x1044, 0x1045, 0x1050, 0x1051, 0x1054, 0x1055,
    0x1100, 0x1101, 0x1104, 0x1105, 0x1110, 0x1111, 0x1114, 0x1115,
    0x1140, 0x1141, 0x1144, 0x1145, 0x1150, 0x1151, 0x1154, 0x1155,
    0x1400, 0x1401, 0x1404, 0x1405, 0x1410, 0x1411, 0x1414, 0x1415,
    0x1440, 0x1441, 0x1444, 0x1445, 0x1450, 0x1451, 0x1454, 0x1455,
    0x1500, 0x1501, 0x1504, 0x1505, 0x1510, 0x1511, 0x1514, 0x1515,
    0x1540, 0x1541, 0x1544, 0x1545, 0x1550, 0x1551, 0x1554, 0x1555,
    0x4000, 0x4001, 0x4004, 0x4005, 0x4010, 0x4011, 0x4014, 0x4015,
    0x4040, 0x4041, 0x4044, 0x4045, 0x4050, 0x4051, 0x4054, 0x4055,
    0x4100, 0x4101, 0x4104, 0x4105, 0x4110, 0x4111, 0x4114, 0x4115,
    0x4140, 0x4141, 0x4144, 0x4145, 0x4150, 0x4151, 0x4154, 0x4155,
    0x4400, 0x4401, 0x4404, 0x4405, 0x4410, 0x4411, 0x4414, 0x4415,
    0x4440, 0x4441, 0x4444, 0x4445, 0x4450, 0x4451, 0x4454, 0x4455,
    0x4500, 0x4501, 0x4504, 0x4505, 0x4510, 0x4511, 0x4514, 0x4515,
    0x4540, 0x4541, 0x4544, 0x4545, 0x4550, 0x4551, 0x4554, 0x4555,
    0x5000, 0x5001, 0x5004, 0x5005, 0x5010, 0x5011, 0x5014, 0x5015,
    0x5040, 0x5041, 0x5044, 0x5045, 0x5050, 0x5051, 0x5054, 0x5055,
    0x5100, 0x5101, 0x5104, 0x5105, 0x5110, 0x5111, 0x5114, 0x5115,
    0x5140, 0x5141, 0x5144, 0x5145, 0x5150, 0x5151, 0x5154, 0x5155,
    0x5400, 0x5401, 0x5404, 0x5405, 0x5410, 0x5411, 0x5414, 0x5415,
    0x5440, 0x5441, 0x5444, 0x5445, 0x5450, 0x5451, 0x5454, 0x5455,
    0x5500, 0x5501, 0x5504, 0x5505, 0x5510, 0x5511, 0x5514, 0x5515,
    0x5540, 0x5541, 0x5544, 0x5545, 0x5550, 0x5551, 0x5554, 0x5555
};

```

Fig. 2. Vector of 256 dilated short integers, each of which 2-dilates the byte that indexes it for Algorithm 1. For a text version of this code, please see the Computer Society Digital Library at <http://doi.ieeecomputersociety.org/10.1109/TC.2007.70814>.

again notationally, $\overleftarrow{j} = 2\overrightarrow{j} = 0xAA00$ (see Fig. 1 for an example of their use as row and column indices).

Notation 2.1. Let w be the number of bits in a (short) word.

Notation 2.2. Let q_ℓ be a quaternary digit.

Notation 2.3. Let o_ℓ be an octal digit.

Each q_ℓ can alternatively be expressed as $q_\ell = 2i_\ell + j_\ell$, where i_ℓ and j_ℓ are bits. Each o_ℓ can alternatively be expressed as $o_\ell = 4i_\ell + 2j_\ell + k_\ell$, where i_ℓ , j_ℓ , and k_ℓ are bits.

Cartesian indices are limited here to w or fewer bits. Morton indices and 2-dilated integers have $2w$ bits. The restriction $w = 16$ can be relaxed for 64-bit architectures, although, frankly, dense matrices do not often have orders exceeding the 65,535 that it provides. If necessary, these algorithms are easily generalized to larger w .

Definition 2.1 [2]. The root of a d -dimensional array has Morton-order index 0. A subarray (block) at Morton-order index i is either an element (scalar) or is composed of 2^d subarrays, with indices $2^d i + 0, 2^d i + 1, \dots, 2^d i + (2^d - 1)$ at the next level.

For $d = 2$, the indexing of a 16×16 matrix is illustrated in Fig. 1. At each level of the recurrence, the four quadrants are labeled *northwest*, *southwest*, *northeast*, and *southeast*, respectively. This yields the \mathbb{N} order to meet the conventions of scientific computation, where matrices are tall rather than wide. Computer graphics, where images are wider than they are tall, uses the transposed \mathbb{Z} order available by exchanging *southwest* and *northeast* just above and using \overleftarrow{i} and \overrightarrow{j} to index rows and columns.

Theorem 2.1 [1]. The Morton index into a matrix is

$$\sum_{\ell=0}^{w-1} q_\ell 4^\ell = 2 \sum_{\ell=0}^{w-1} i_\ell 4^\ell + \sum_{\ell=0}^{w-1} j_\ell 4^\ell \text{ and corresponds to the Cartesian index for row } i = \sum_{\ell=0}^{w-1} i_\ell 2^\ell \text{ and column } j = \sum_{\ell=0}^{w-1} j_\ell 2^\ell.$$

The concatenated quaternary digits above index the nested quadrants of a matrix. The octal digits, below, index nested octants of a 3D array.

Theorem 2.2 [2]. The Morton index into a 3D array is

$$\sum_{\ell=0}^{w-1} o_\ell 8^\ell = 4 \sum_{\ell=0}^{w-1} i_\ell 8^\ell + 2 \sum_{\ell=0}^{w-1} j_\ell 8^\ell + \sum_{\ell=0}^{w-1} k_\ell 8^\ell \text{ and corresponds to the Cartesian index for row } i = \sum_{\ell=0}^{w-1} i_\ell 2^\ell, \text{ column } j = \sum_{\ell=0}^{w-1} j_\ell 2^\ell, \text{ and rod } k = \sum_{\ell=0}^{w-1} k_\ell 2^\ell.$$

These theorems reflect the so-called “bit interleaving” [15], [16, pp. 53-55], [1], also described in [17]. The purpose of the following algorithms is to convert between the binary representations of $i = \sum_{\ell=0}^{w-1} i_\ell 2^\ell$ and of $\overrightarrow{i} = \sum_{\ell=0}^{w-1} i_\ell 2^{d\ell}$, that is, the dilation of integers.

The following two sections define four functions each for casting ordinary integers to and from 2-dilated and 3-dilated integers by using table lookup and word width arithmetic. The former functions run in about $w/8$ steps (on tables of 2^8 entries), but the latter use roughly $\lg w$ steps. Even with w growing to 32, addressing a square matrix of order $4 \cdot 10^9$, these proportions are about the same. Section 4.4 offers two more functions for even wider dilations.

3 CONVERTING VIA TABLE LOOKUP

The following algorithms are conversions and inverse conversions (*dilations* and *contractions*) from integers to 2-dilations: $\sum_{\ell=0}^{w-1} i_\ell 2^\ell \Leftrightarrow \sum_{\ell=0}^{w-1} i_\ell 4^\ell$ with inline instructions. The first casting conversions use table lookup. They require a vector of 256 entries, either short integers or bytes. Its aligned footprint in an L1 cache is very small and the inline code is short.

For the first pair of algorithms, assume that $d = 2$, thus spreading bits apart by one, that the dilated integers are 32 bits

```

const unsigned char undilate_tab2[256] = {
    0x00,0x01,0x10,0x11, 0x02,0x03,0x12,0x13, 0x20,0x21,0x30,0x31, 0x22,0x23,0x32,0x33,
    0x04,0x05,0x14,0x15, 0x06,0x07,0x16,0x17, 0x24,0x25,0x34,0x35, 0x26,0x27,0x36,0x37,
    0x40,0x41,0x50,0x51, 0x42,0x43,0x52,0x53, 0x60,0x61,0x70,0x71, 0x62,0x63,0x72,0x73,
    0x44,0x45,0x54,0x55, 0x46,0x47,0x56,0x57, 0x64,0x65,0x74,0x75, 0x66,0x67,0x76,0x77,
    0x08,0x09,0x18,0x19, 0x0A,0x0B,0x1A,0x1B, 0x28,0x29,0x38,0x39, 0x2A,0x2B,0x3A,0x3B,
    0x0C,0x0D,0x1C,0x1D, 0x0E,0x0F,0x1E,0x1F, 0x2C,0x2D,0x3C,0x3D, 0x2E,0x2F,0x3E,0x3F,
    0x48,0x49,0x58,0x59, 0x4A,0x4B,0x5A,0x5B, 0x68,0x69,0x78,0x79, 0x6A,0x6B,0x7A,0x7B,
    0x4C,0x4D,0x5C,0x5D, 0x4E,0x4F,0x5E,0x5F, 0x6C,0x6D,0x7C,0x7D, 0x6E,0x6F,0x7E,0x7F,
    0x80,0x81,0x90,0x91, 0x82,0x83,0x92,0x93, 0xA0,0xA1,0xB0,0xB1, 0xA2,0xA3,0xB2,0xB3,
    0x84,0x85,0x94,0x95, 0x86,0x87,0x96,0x97, 0xA4,0xA5,0xB4,0xB5, 0xA6,0xA7,0xB6,0xB7,
    0xC0,0xC1,0xD0,0xD1, 0xC2,0xC3,0xD2,0xD3, 0xE0,0xE1,0xF0,0xF1, 0xE2,0xE3,0xF2,0xF3,
    0xC4,0xC5,0xD4,0xD5, 0xC6,0xC7,0xD6,0xD7, 0xE4,0xE5,0xF4,0xF5, 0xE6,0xE7,0xF6,0xF7,
    0x88,0x89,0x98,0x99, 0x8A,0x8B,0x9A,0x9B, 0xA8,0xA9,0xB8,0xB9, 0xAA,0xAB,0xBA,0xBB,
    0x8C,0x8D,0x9C,0x9D, 0x8E,0x8F,0x9E,0x9F, 0xAC,0xAD,0xBC,0xBD, 0xAE,0xAF,0xBE,0xBF,
    0xC8,0xC9,0xD8,0xD9, 0xCA,0xCB,0xDA,0xDB, 0xEA,0xEB,0xFA,0xFB, 0xEC,0xED,0xFC,0xFD, 0xEE,0xEF,0xFE,0xFF,
    0xCC,0xCD,0xDC,0xDD, 0xCE,0xCF,0xDE,0xDF, 0xEC,0xED,0xFC,0xFD, 0xEE,0xEF,0xFE,0xFF
};

```

Fig. 3. Vector of all possible bytes, each indexed by the folding of its respective 2-dilated integer for Algorithm 2. For a text version of this code, please see the Computer Society Digital Library at <http://doi.ieeecomputersociety.org/10.1109/TC.2007.70814>.

```

const unsigned char dilate_tab3[256] = {
    0x00,0x01,0x08,0x09, 0x40,0x41,0x48,0x49, 0x02,0x03,0x0A,0x0B, 0x42,0x43,0x4A,0x4B,
    0x10,0x11,0x18,0x19, 0x50,0x51,0x58,0x59, 0x12,0x13,0x1A,0x1B, 0x52,0x53,0x5A,0x5B,
    0x80,0x81,0x88,0x89, 0xC0,0xC1,0xC8,0xC9, 0x82,0x83,0x8A,0x8B, 0xC2,0xC3,0xCA,0xCB,
    0x90,0x91,0x98,0x99, 0xD0,0xD1,0xD8,0xD9, 0x92,0x93,0x9A,0x9B, 0xD2,0xD3,0xDA,0xDB,
    0x04,0x05,0x0C,0x0D, 0x44,0x45,0x4C,0x4D, 0x06,0x07,0x0E,0x0F, 0x46,0x47,0x4E,0x4F,
    0x14,0x15,0x1C,0x1D, 0x54,0x55,0x5C,0x5D, 0x16,0x17,0x1E,0x1F, 0x56,0x57,0x5E,0x5F,
    0x84,0x85,0x8C,0x8D, 0xC4,0xC5,0xCC,0xCD, 0x86,0x87,0x8E,0x8F, 0xC6,0xC7,0xCE,0xCF,
    0x94,0x95,0x9C,0x9D, 0xD4,0xD5,0xDC,0xDD, 0x96,0x97,0x9E,0x9F, 0xD6,0xD7,0xDE,0xDF,
    0x20,0x21,0x28,0x29, 0x60,0x61,0x68,0x69, 0x22,0x23,0x2A,0x2B, 0x62,0x63,0x6A,0x6B,
    0x30,0x31,0x38,0x39, 0x70,0x71,0x78,0x79, 0x32,0x33,0x3A,0x3B, 0x72,0x73,0x7A,0x7B,
    0xA0,0xA1,0xA8,0xA9, 0xE0,0xE1,0xE8,0xE9, 0xA2,0xA3,0xAA,0xAB, 0xE2,0xE3,0xEA,0xEB,
    0xB0,0xB1,0xB8,0xB9, 0xF0,0xF1,0xF8,0xF9, 0xB2,0xB3,0xBA,0xBB, 0xF2,0xF3,0xFA,0xFB,
    0x24,0x25,0x2C,0x2D, 0x64,0x65,0x6C,0x6D, 0x26,0x27,0x2E,0x2F, 0x66,0x67,0x6E,0x6F,
    0x34,0x35,0x3C,0x3D, 0x74,0x75,0x7C,0x7D, 0x36,0x37,0x3E,0x3F, 0x76,0x77,0x7E,0x7F,
    0xA4,0xA5,0xAC,0xAD, 0xE4,0xE5,0xEC,0xED, 0xA6,0xA7,0xAE,0xAF, 0xE6,0xE7,0xEE,0xEF,
    0xB4,0xB5,0xBC,0xBD, 0xF4,0xF5,0xFC,0xFD, 0xB6,0xB7,0xBE,0xBF, 0xF6,0xF7,0xFE,0xFF
};

```

Fig. 4. Vector of all 256 bytes, each of which 3-dilates the byte that indexes it *and* each indexed by the folding of its respective 3-dilated integer for *both* Algorithms 3 and 4. For a text version of this code, please see the Computer Society Digital Library at <http://doi.ieeecomputersociety.org/10.1109/TC.2007.70814>.

wide, that ordinary integer indices are 16 bits wide, and that all integers are unsigned. Their generalization is straightforward for 64-bit indexing into large arrays. Generalizations to $d = 3$ follow.

The table for 2-dilation of Algorithm 1 appears as Fig. 2 [18], [19]. It is composed of 256 short integers, occupying twice the necessary bits but avoiding final operations to mask out alternating bits from them. Those extra operations are illustrated in Algorithm 3 for 3-dilation.

Algorithm 1

```

inline unsigned int dilate_2(unsigned short x){
    return dilate_tab2[ 0xFF & x ]
    | (dilate_tab2[ (0xFFFF & x) >>8 ] <<16);
}

```

Algorithm 2

```

inline unsigned short undilate_2(unsigned int x){
    return undilate_tab2[0xFF & ((x>>7) | x) ]
    | (undilate_tab2[0xFF & (((x>>7) | x) >>16)]
    << 8);
}

```

To invert a dilation, the dilated integer is split and folded into a single pattern, byte by byte, which then indexes a table (Fig. 3) to recover its contraction. This folding halves the fetches from others' tables [18], [19]. Algorithms 1-4 use bitwise inclusive-OR $|$ to sum bit patterns. Alternatives are simple addition $+$ and bitwise exclusive-OR \wedge , because the patterns are bitwise disjoint.

Algorithms 3 and 4 use the precomputed vector of 256 bytes in Fig. 4, instead of two: one of shorts and another of bytes. The

algorithms are written for 3-dilations of 16-bit short integers to 32-bit integers and vice versa.

Algorithm 3

```

inline unsigned int dilate_3(unsigned short x){
    return ((
        ((dilate_tab3[(0xFFFF & x) >>8] ) <<24)
        | dilate_tab3[ 0xFF & x ]
        ) * 0x010101) & 0x49249249;
}

```

Algorithm 4

```

inline unsigned short undilate_3(unsigned int x){
    return dilate_tab3[0xFF & (((x>>8) | x) >>8)
    | (dilate_tab3[0xFF & (((x>>8) | x) >>8)
    | x) >>24]] <<8);
}

```

Just as Algorithm 2 halves the space of Algorithm 1 (Fig. 3 versus Fig. 2), the next theorem nearly halves their total space once again.

Theorem 3.1. *The tables of bytes required for Algorithms 3 and 4 coincide.*

Proof [suggested by a referee]. A table entry effects a dilating/folding permutation P of the bit at an individual Position i , $0 \leq i < 8$, to Position $(3 \cdot i) \bmod 8$ and it thus permutes all eight bits in a single lookup. This permutation is its own inverse because $P(P(i)) = (3 \cdot 3 \cdot i) \bmod 8 = i \bmod 8 = i$. \square

```

inline u_short undilate_2(u_int t) {
    t = (t | (t >> 1)) & 0x33333333;
    t = (t | (t >> 2)) & 0x0F0F0F0F;
    t = (t | (t >> 4)) & 0x00FF00FF;
    t = (t | (t >> 8)) & 0x0000FFFF;
    return((u_short) t);
}

```

Algorithm 5:

```

inline u_int dilate_2(u_short t) {
    u_int r = t;
    r = (r | (r << 8)) & 0x00FF00FF;
    r = (r | (r << 4)) & 0x0F0F0F0F;
    r = (r | (r << 2)) & 0x33333333;
    r = (r | (r << 1)) & 0x55555555;
    return(r);
}

```

Fig. 5. Inline algorithms described by Merkey for 2-dilation and 2-undilation. `u_int` and `u_short` are assumed to be typedefed to `unsigned int` and `unsigned short`, respectively [20], [19]. The left one remains unnumbered because it is replaced by Algorithm 6. For a text version of this code, please see the Computer Society Digital Library at <http://doi.ieeecomputersociety.org/10.1109/TC.2007.70814>.

A similar proof applies for any d -dilation and permuting table of 2^p different p -bit integers such that $d^2 \bmod p = 1$. For example, one table of 256 bytes also suffices for 5-dilation and contraction, although computations like the ones in the next sections are faster.

4 CONVERSION WITHOUT TABLE LOOKUP

Conversion can also be done entirely within the processor without the cache impact of table lookup. A review of existing approaches follows, based on shifts and masking, which leads to new ones based on integer multiplication. This section concludes by considering the abstract case for all d . Discussed first are d -undilation algorithms for $d = 2, 3$.

As before, consider the case where the dilated integer is contained in a w -bit (32-bit) word. The w limit constrains the number of significant digits in the undilated integer to $s = \lfloor w/d \rfloor$. As shown in several cases in the following, abstract multiplication generates high-order overflow bits that are ignored.

Notation 4.1. Let s denote the number of significant dilatable bits in an undilated integer.

4.1 Shift-Or Algorithms

Merkey [20] presents straight-line codes, based on an algorithm by Stocco and Schrack [19], which use repeated shift-OR operations and masks for conversion between normal and dilated integers. They can also be found without citation among sample codes in machine manuals [21, pp. 136-193, 184-185]. Instances of the shift-based 2-undilation and 2-dilation algorithms for $s = 16$ are shown in Fig. 5. Each algorithm works in four rounds, where a round is comprised of three operations: a shift and two bitwise operations. Merkey presents a variant that performs one shift and three bitwise operations per round, but this works equally well with superscalar instructions. Another round is necessary for $w = 64$ and $s = 32$.

Fig. 5 illustrates the whole family of algorithms. In `undilate_2`, for instance, each round coalesces pairs of groups from the previous round, that is, that algorithm performs “binary coalescing.” Unnumbered in this figure because it is replaced in the following, it uses a mask to remove the redundant “middle muddle” that results from reflexive bit merges. It collapses the dilation in only $\lceil \lg s \rceil$ rounds.

4.2 Undilation via Multiplication

With fast integer multiplication, the shift-ORs in a round of the undilation algorithm might be reduced to a single multiplication that implements them. When a constant factor has only a few bits set, its multiplication might be better replaced by inline shift-adds at compile time or (even if a variable) decoded by the processor to shift-adds at runtime. For example, multiplications by integers like 3 or 17 that have only 2 bits set in their binary representation can be implemented by a single shift-add. In writing such a multiplication, we intend it to allow such an implementation, if faster.

Indeed, we have seen compilers that enforce it (sometimes *slowing* the code). Faster multipliers or relatively slower barrel shifters might favor the multiply as coded.

The multiplication is safe here because the intervening/dilating 0 bits absorb any carries from the implied addition within the middle muddle, which is about to be masked off, anyway. Since any dilated integer has a sparse set of significant bit positions, it is safe to implement the shift-ORs of similar undilations by using multiplication. For instance, 2 and 3-undilation can be implemented as follows:

Algorithm 6

```

inline u_short undilate_2(u_int t){
    t = (t * 3) & 0x66666666;
    t = (t * 5) & 0x78787878;
    t = (t * 17) & 0x7F807F80;
    t = (t * 257) & 0x7FFF8000;
    return ((u_short) (t >> 15));
}

```

Algorithm 7

```

inline u_short undilate_3(u_int t){
    t = (t * 0x00015) & 0x0E070381;
    t = (t * 0x01041) & 0x0FF80001;
    t = (t * 0x40001) & 0x0FFC0000;
    return ((u_short) (t >> 18));
}

```

The moderately experienced programmer will recognize the constant factors in Algorithm 6 as successors of powers of powers of 2 (Fermat primes even) and, so, we express them as decimal numbers. In fact, and as derived in Section 4.4, in Round i , the factor is composed of 2 one bits, separated by $\frac{1}{2}2^i - 1$ zero bits. In Algorithm 7, they are presented as hexadecimal numbers, revealing their similarly regular bit patterns: one bits are separated by $\frac{2}{3}3^i - 1$ zero bits in Round i .

Algorithm 6 mimics Fig. 5’s `undilate_2`, with the difference that the partial results shift to the left as rounds proceed because the products emulate right shifts. For example, its first statement is mimicked almost directly by the assignment `t = (t * 3) & 0x66666666` in Algorithm 6, except for that shift. Because of it, a cumulative correcting shift is necessary in the return value.

In summary, Algorithm 6 requires nine operations and Algorithm 7 requires seven. For 64-bit dilated indices, an extra round is needed by both algorithms, raising the instruction counts to 11 and 9, respectively. See Table 1 for comparisons.

TABLE 1

Instruction Counts for Undilation and Dilation for Practical Algorithms Using Table Lookup, Shift-Based, and Multiplication-Based Algorithms

Strategy	Alg.	Dilation			Undilation		
		fetch	shift	mask	fetch	shift	mask
Table Lookup							
$d = 2$:	1,2	2	2	3	2	2	4
$d = 3$:	3,4	2	2 ¹	3 ¹	2	4	5
Processor-local		multiply	shift	mask	multiply	shift	mask
$d = 2$:	5,6	0	4	8	4	1	4
$d = 3$:	8,7	4	0	4	3	1	3

¹ Multiplication by 0×010101 in Algorithm 3 is not counted here. Depending on your processor, it may be faster as the constant multiplication or as two shifts and two additions.

4.3 Dilation

Algorithm 5 stands as our processor-local algorithm for 2-dilation. Multiplication alone cannot replace its shift-OR operations because it introduces carry bits whose effects survive the subsequent masking. For example, one is tempted to substitute $t = t * 257 \& \text{mask}$ for its first statement $t = (t|(t << 8)) \& \text{mask}$, but that fails, for example, at $t = 257$, because of the carry that shifts into the 2^9 bit of the product. Indeed, we are not aware of any multiplication-based approach to 2-dilation that uses fewer instructions than the shift-based approach.

The multiplication-based approach, however, *can* be used to 3-dilate an integer via “binary splitting” with fewer instructions than a shift-based binary-splitting algorithm.

Algorithm 8

```

inline u_int dilate_3(u_short t) {
    u_int r = t;
    r = (r * 0x10001) & 0xFF0000FF;
    r = (r * 0x00101) & 0x0F00F00F;
    r = (r * 0x00011) & 0xC30C30C3;
    r = (r * 0x00005) & 0x49249249;
    return(r);
}

```

As before, the constant factors in Algorithm 8 follow a pattern that can be inferred from their hexadecimal presentation: one bits are separated by $2^{5-i} - 1$ zero bits at Round i . This algorithm delivers a binary splitting pattern of partial results: At the end of each round, there is a maximum group size which changes as 8, 4, 2, and then 1.

While Algorithm 8 needs as many rounds as a shift-based approach, each round requires just two instructions, multiply and mask, that might become the shift-based algorithms’ three. Thus, dilating for $d = 3$ and $s = 10$ takes 8 instructions rather than 12 and, for 64-bit dilated integers, it would require only 10.

Although a “ternary split” could perform 3-dilation in fewer rounds, such an approach is inferior for both the multiplication-based and shift-based approaches. Multiplication fails because of the presence of undesirable carries. As before, shift-based approaches would also have an excessive number of operations per round, more than nullifying any reduction in the number of rounds. However, as shown in the next section, multiplication implements ternary split well for 4-dilation and, in general, the $(d - 1)$ -ary split does well for wider d -dilations.

4.4 Even Wider Dilations

The case of arbitrary $d > 1$ and $s > 1$ is now addressed. To be specific, how can one derive the “magic” constants above and how is one convinced of the correctness of these algorithms? (Those constants are denoted b, c, y, z in the following, subscripted for dilation d and the round i .) In the arbitrary case, s is subject only to

the condition that $1 < s \leq \lfloor w/d \rfloor$, where w denotes the natural word width of the processor. The general algorithms follow.

Algorithm 9

DILATE(t, s, d) (where $d > 2$)

- 1) $r \leftarrow t$;
- 2) for $i = 1, 2, \dots, \lceil \log_{d-1} s \rceil$ do
 $r \leftarrow (r \times b_{d,i}) \text{ AND } y_{d,i}$.
- 3) return r .

Algorithm 10

UNDILATE(t, s, d)

- 1) $r \leftarrow t$;
- 2) for $i = 1, 2, \dots, \lceil \log_d s \rceil$ do
 $r \leftarrow (r \times c_{d,i}) \text{ AND } z_{d,i}$.
- 3) $a \leftarrow d(s - 1) + 1$;
- 4) return $\lfloor r/2^{a-s} \rfloor$.

When $d > 2$, the multiplication-based algorithms use $2\lceil \log_{d-1} s \rceil$ instructions for d -dilation and, for any $d \geq 2$, $2\lceil \log_d s \rceil + 1$ instructions for d -undilation. The bases of these logarithms lead to the terms $(d - 1)$ -recursion and d -recursion. Compared with the shift-based algorithms’ bound of $3\lceil \log_2 s \rceil$ instructions, the multiplication-based algorithms use far fewer instructions, particularly when d becomes large due to the linearity of d in the base of the logarithm. (Table 1 presents exact counts for practical values of d and s .)

Consider now the correctness of the algorithms and the values of b, c, y , and z .

Notation 4.2. Let $x_{p,q}$ denote $\sum_{\ell=0}^{p-1} 2^{\ell q}$ for integers $p, q \geq 1$.

The binary expansion of $x_{p,q}$ has p ones, with each pair of successive ones separated by exactly $q - 1$ zeros and with the rightmost bit equal to 1.

The correctness of the undilation Algorithm 10 is implied by an induction on two invariants. They determine what the constants $c_{d,i}$ and $z_{d,i}$ should be. The initial/base case is used in the following as if it occurred at “the end of” a null Round 0:

1. For all rounds, the leftmost bit of r stays in the same bit position, that is, position $d(s - 1) + 1$.
2. At the end of Round i for $i = 0, 1, \dots$, there are groups of consecutive significant digits of size d^i , with the rightmost group possibly smaller if s is not a power of d . Each group is separated from the one before it by $d^i(d - 1)$ zero bits.

It is easy to verify that the multiplication constant for the $(i + 1)$ th round $c_{d,i+1}$ should be $x_{d,(d-1)d^i}$. This constant can be more than w bits long, but, by the properties of modular arithmetic, only the low-order w bits of any constant are necessary for the result to be correct. For instance, when $d = 3$, $s = 10$, and $w = 32$, $c_{3,3} = x_{3,18} = 2^{36} + 2^{18} + 1 = 0 \times 1000040001$, a 37-bit value. However, the constant 0×40001 in Algorithm 7 is just the low-order 32 bits of $c_{3,3}$.

Algorithm 9, d -dilation for $d > 2$ in $t = \lceil \log_{d-1} s \rceil$ rounds, has invariants.

3. For all rounds, the rightmost bit of r stays in Position 0.
4. At the end of Round i for $i = 0, 1, \dots, t$, there are groups of consecutive significant digits of size $(d - 1)^{t-i}$ plus possibly one (the leftmost) group of smaller size if s is not a power of $(d - 1)$. Each group is separated from the one before it by $(d - 1)^{t-i+1}$ blank bits.

Therefore, one round of d -dilation reduces the maximum size of a group of consecutive significant digits by a factor of $(d - 1)$, which may be viewed as a $(d - 1)$ -way split. Fig. 6 illustrates why this




Fig. 6. Schematic view of round $i + 1$ in the general dilation algorithm. Each rectangle represents a block of consecutive significant digits. The top row is the state at the end of round i , the bottom row is the state at the end of round $i + 1$, and the middle rows show the effect of the multiplication in round $i + 1$. The dotted lines indicate the parts of the product that are retained by the mask operations.

works. As was shown, one can use a multiplication to create $(d - 1)$ adjacent copies of a block of $(d - 1)^j$ contiguous significant bits. From these adjacent copies, one can extract the appropriate contiguous blocks of size $(d - 1)^{j-1}$, which are spaced exactly $(d - 1)^j$ bits apart, as required by Invariant 4. The multiplication constants are given by $b_{d,i} = x_{d,(d-1)^{i-1}}$.

5 TIMINGS AND CONCLUSION

Algorithms 1-8 have been tested on PowerPCs G4 and G5, Pentiums 2 and 4, Xeon, an Itanium, Athlons, Opterons, and an IBM POWER5. Timings were taken for repeated runs with the small tables of Algorithms 1-4 preloaded in cache, a reasonable assumption when they are being used often, as in a loop. (The alternative of condemning the cache after every test to simulate rare random conversions is too stringent.) All are very fast. Algorithms 1 and 2 are consistently faster than Algorithms 5 and 6. Algorithms 3 and 4 compare differently to 8 and 7 on different machines. For wider dilations ($d > 3$), unpacking the table entries becomes tedious and the direct multiplication-based algorithms shine.

TABLE 2
Average Times for the Algorithms Measured
in Processor Cycles on Different Machines

Processor		$d = 2$ for Matrices			
		Table Lookup		Multiply	
		Dilate Alg. 1	Undilate Alg. 2	Dilate Alg. 5	Undilate Alg. 6
Itanium	gcc 3.4	2.64	4.65	9.66	4.65
Itanium	icc 8.1	3.11	5.11	10.12	8.12
Opteron	gcc 4.1.1	1.05	2.57	6.08	4.54
Pentium4	gcc 3.4.6	4.04	6.55	15.17	10.85
Pentium-M	gcc 3.3.6	3.01	4.14	8.37	7.96
POWER5	gcc 3.3.3	1.08	2.95	8.39	4.85
Xeon	gcc 3.4.5	4.06	6.48	15.18	10.45
Measured in average processor cycles.		$d = 3$			
		Table Lookup		Multiply	
		Dilate Alg. 3	Undilate Alg. 4	Dilate Alg. 8	Undilate Alg. 7
Itanium	gcc 3.4	7.68	6.65	7.65	6.65
Itanium	icc 8.1	10.12	8.12	8.12	8.12
Opteron	gcc 4.1.1	3.10	5.05	5.19	2.99
Pentium4	gcc 3.4.6	10.81	8.56	14.16	13.72
Pentium-M	gcc 3.3.6	5.02	4.83	8.61	5.56
POWER5	gcc 3.3.3	4.88	4.97	8.64	2.98
Xeon	gcc 3.4.5	10.71	8.49	13.94	13.68

Table 1 compares the operation counts for the eight principal algorithms. The trade-off between fetch and multiplication is apparent. Table 2 compares the runtimes measured in clock cycles for these algorithms, run on several machines. The indicated C compilers and versions were used only with option -O3.

This paper presents several results that advance the use of dilated integers for indexing Morton-ordered matrices. First, it is an introduction to those structures and how simply they deliver block locality with conventional Cartesian indexing. Second, it reintroduces efficient processor-based algorithms that should be more widely known and establishes new time-efficient and space-efficient table-based alternatives. Of course, either could be displaced by hardwired bit shufflers. Third, it generalizes the former algorithms with d -ary and $(d - 1)$ -ary recurrences for the general case.

The expansion of the memory hierarchy and the advent of chip multiprocessors create a need for arrays with highly local access within any cache-resident block. These algorithms allow dilated integers and methods for direct arithmetic on them [12] and help deliver that locality with ordinary Cartesian indexing on ordinary hardware.

ACKNOWLEDGMENTS

The authors thank the sponsors of Schloss Dagstuhl in the Saarland, for creating the magnificent environment and international interactions that inspired this paper. Thanks also to the students in the B679 Seminar at Indiana University, particularly K. Patrick Lorton, who timed these codes on their favorite machines, and to the anonymous diligent referees. This work was supported in part by the US National Science Foundation under grants CCR-0073491, EIA-0202048, and CCF-0541364. POWER5 tests used IU computing systems.

REFERENCES

- [1] G.M. Morton, "A Computer-Oriented Geodetic Data Base and a New Technique in File Sequencing," technical report, IBM Ltd., Ottawa, Mar. 1966.
- [2] D.S. Wise, "Ahnentafel Indexing into Morton-Ordered Arrays, or Matrix Locality for Free," *Euro-Par 2000—Parallel Processing*, A. Bode, T. Ludwig, W. Karl, and R. Wismüller, eds., pp. 774-883, Springer, 2000. http://dx.doi.org/10.1007/3-540-44520-X_108
- [3] J.D. Frens and D.S. Wise, "QR Factorization with Morton-Ordered Quadtree Matrices for Memory Reuse and Parallelism," *Proc. Ninth ACM SIGPLAN Symp. Principles and Practice of Parallel Program, SIGPLAN Not.* vol. 38, no. 10, pp. 144-154, Oct. 2003. <http://dx.doi.org/10.1145/781498.781525>

- [4] J. Sang Park, M. Penner, and V.K. Prasanna, "Optimizing Graph Algorithms for Improved Cache Performance," *IEEE Trans. Parallel Distrib. Syst.*, vol. 15, no. 9, pp. 769-782, Sept. 2004.
<http://dx.doi.org/10.1109/TPDS.2004.44>
- [5] M. Frigo, C.E. Leiserson, H. Prokop, and S. Ramachandran, "Cache-Oblivious Algorithms," *Proc. 40th Ann. Symp. Foundations of Computer Science*, pp. 285-298, Oct. 1999.
<http://dx.doi.org/10.1109/SFFCS.1999.814600>
- [6] M.D. Adams and D.S. Wise, "Seven at One Stroke: Results from a Cache-Oblivious Paradigm for Scalable Matrix Algorithms," *Proc. 2006 Workshop Memory System Performance and Correctness*, pp. 41-50, Oct. 2006.
<http://dx.doi.org/10.1145/1178597.1178604>
- [7] J.G. Siek and A. Lumsdaine, "The Matrix Template Library: Generic Components for High-Performance Scientific Computing," *Computing in Science and Eng.*, vol. 1, no. 6, pp. 70-78, Nov. 1999.
<http://dx.doi.org/10.1109/5992.805137>
- [8] W.D. Clinger, "How to Read Floating-Point Numbers Accurately," *SIGPLAN Not.*, vol. 39, no. 4, best of PLDI, pp. 360-371, Apr. 2004, originally published vol. 25, no. 6, pp. 92-101, June 1990.
<http://dx.doi.org/10.1145/989393.989430>
- [9] G.L. Steele Jr. and J.L. White, "How to Print Floating-Point Numbers Accurately," *SIGPLAN Not.*, vol. 39, no. 4, best of PLDI, pp. 372-389, Apr. 2004, originally published vol. 25, no. 6, pp. 112-126, June 1990.
<http://dx.doi.org/10.1145/989393.989431>
- [10] M. Thottethodi, S. Chatterjee, and A.R. Lebeck, "Tuning Strassen's Matrix Multiplication for Memory Efficiency," *Proc. Supercomputing*, chapter 36, Nov. 1998. <http://dx.doi.org/10.1109/SC.1998.10045>
- [11] V. Valsalam and A. Skjellum, "A Framework for High-Performance Matrix Multiplication Based on Hierarchical Abstractions, Algorithms and Optimized Low-Level Kernels," *Concur. Comp. Prac. Exper.*, vol. 14, no. 10, pp. 805-839, 2002. <http://dx.doi.org/10.1002/cpe.630>
- [12] G. Schrack, "Finding Neighbors of Equal Size in Linear Quadtrees and Octrees in Constant Time," *CVGIP: Image Underst.*, vol. 55, no. 3, pp. 221-230, May 1992.
- [13] M.D. Adams and D.S. Wise, "Fast Additions on Masked Integers," *SIGPLAN Not.*, vol. 41, no. 5, pp. 39-45, May 2006.
<http://dx.doi.org/10.1145/1149982.1149987>
- [14] D.E. Knuth, *Fundamental Algorithms*, third ed., vol. 1, series the Art of Computer Programming. Addison-Wesley, 1997.
- [15] G. Peano, "Sur une Courbe, Qui Remplit Toute une Aire Plaine," *Math. Ann.*, vol. 36, pp. 157-160, 1890.
- [16] K.D. Tocher, "The Application of Automatic Computers to Sampling Experiments," *J. Roy. Statist. Soc. Ser. B*, vol. 16, no. 1, pp. 39-61, 1954.
- [17] H. Samet, *The Design and Analysis of Spatial Data Structures*, section 2.7. Addison-Wesley, 1990.
- [18] F.C. Holroyd and D.C. Mason, "Efficient Linear Quadtree Construction Algorithm," *Image Vision Comput.*, vol. 8, no. 3, pp. 218-224, Aug. 1990.
- [19] L. Stocco and G. Schrack, "Integer Dilation and Contraction for Quadtrees and Octrees," *Proc. IEEE Pacific Rim Conf. Comm., Computers, and Signal Processing*, pp. 426-428, May 1995.
<http://dx.doi.org/10.1109/PACRIM.1995.519560>
- [20] P. Merkey, "Z-Ordering and UPC," technical report, Michigan Technological Univ., <http://www.upc.mtu.edu/papers/zorder.pdf>, June 2003.
- [21] AMD Athlon Processor x86 Code Optimization Guide. Advanced Micro Devices, Sunnyvale, Calif., publication 22007, http://www.amd.com/us-en/assets/content_type/white_papers_and_tech_docs/22007.pdf, Feb. 2002.

► For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.